

Building an streaming cloud production data pipeline with Apache Beam

Esteban Maya Cadavid



About me



- Software Engineer and consultant
- VC Investor.
- Co-Organizer of PyCon Colombia Conference <https://pycon.co/>
- Co-Organizer of Python Medellin / Colombia Community
- Born and base in Colombia, Medellin

Agenda

- Introduction.
 - Introduction to Beam.
 - Beam model and concepts.
 - Overview Beam io packages.
 - Beam supported languages.
- Architecture
 - A common ETL flow
 - Streaming data pipeline architecture overview.
 - Tips and tricks
 - Debugging streaming pipeline locally (Demo live).
- Deploy
 - Deploying the data pipeline in cloud runner (Demo Live).
- Q/A session.

Introduction to Apache beam

What is Apache Beam?

- An advanced unified programming model.
- Implement batch and streaming data jobs agnostic that run in cloud, locally and many others runners/execution engines.

Why Beam?

- Write the code one time in (Java, Python, Go) and run anywhere.
- It's an agnostic technology which means you can run it in cloud runners like Google Dataflow. Or you can run it in your own infrastructure/machine using spark, flink, samza.
- It's open source, portable, flexible and extensible.

Beam Model and Concepts

PCollections

The PCollection abstraction represents a potentially distributed, multi-element data set.

For example if you are reading a csv file like this.

1	id,firstname,lastname,email,profession
2	100,Micheline,Hartnett,Micheline.Hartnett@yopmail.com,police officer
3	101,Heddie,Georas,Heddie.Georas@yopmail.com,police officer
4	102,Corene,Genna,Corene.Genna@yopmail.com,developer
5	103,Chickie,Bivins,Chickie.Bivins@yopmail.com,developer
6	104,Sharai,Muriel,Sharai.Muriel@yopmail.com,developer
7	105,Gerianna,Isacco,Gerianna.Isacco@yopmail.com,worker
8	106,Jolyn,Khorma,Jolyn.Khorma@yopmail.com,worker
9	107,Vere,Orelee,Vere.Orelee@yopmail.com,police officer
10	108,Devina,Belanger,Devina.Belanger@yopmail.com,firefighter
11	109,Florie,Helfand,Florie.Helfand@yopmail.com,firefighter

The dataset above is our PCollection.

PTransforms

Transforms are the operations in your pipeline. You provide processing logic in the form of a function object (colloquially referred to as “user code”), and your user code is applied to each element of an input PCollection

For example

```
def get_columns(data):  
    id, firstname, lastname, email, profession = data.split(",")  
  
    return [{  
        "id": id,  
        "firstname": firstname,  
        "lastname": lastname,  
        "email": email,  
        "profession": profession  
    }]
```

The function above is a PTransform which is applied to get the data over each element from our previous PCollection and parse them in a dictionary object.

Beam io packages

Apache beam support most of the know well industry sources and destinations technologies. Here is some of them

Cloud storage (Google storage GCS, AWS S3, Azure blobstorage...).

Queue brokers (Kafka, Google pub/sub, AWS SQS...)

Files (Plain text, csv, parquet...)

Databases (Mongo, JDBC, Snowflake, Bigquery, Cassandra...)

You can consult the entire list according the language here.

- https://beam.apache.org/releases/pydoc/2.32.0/apache_beam.io.html
- https://pkg.go.dev/github.com/apache/beam/sdks/go/pkg/beam?utm_source=godoc
- <https://beam.apache.org/releases/javadoc/2.32.0/>

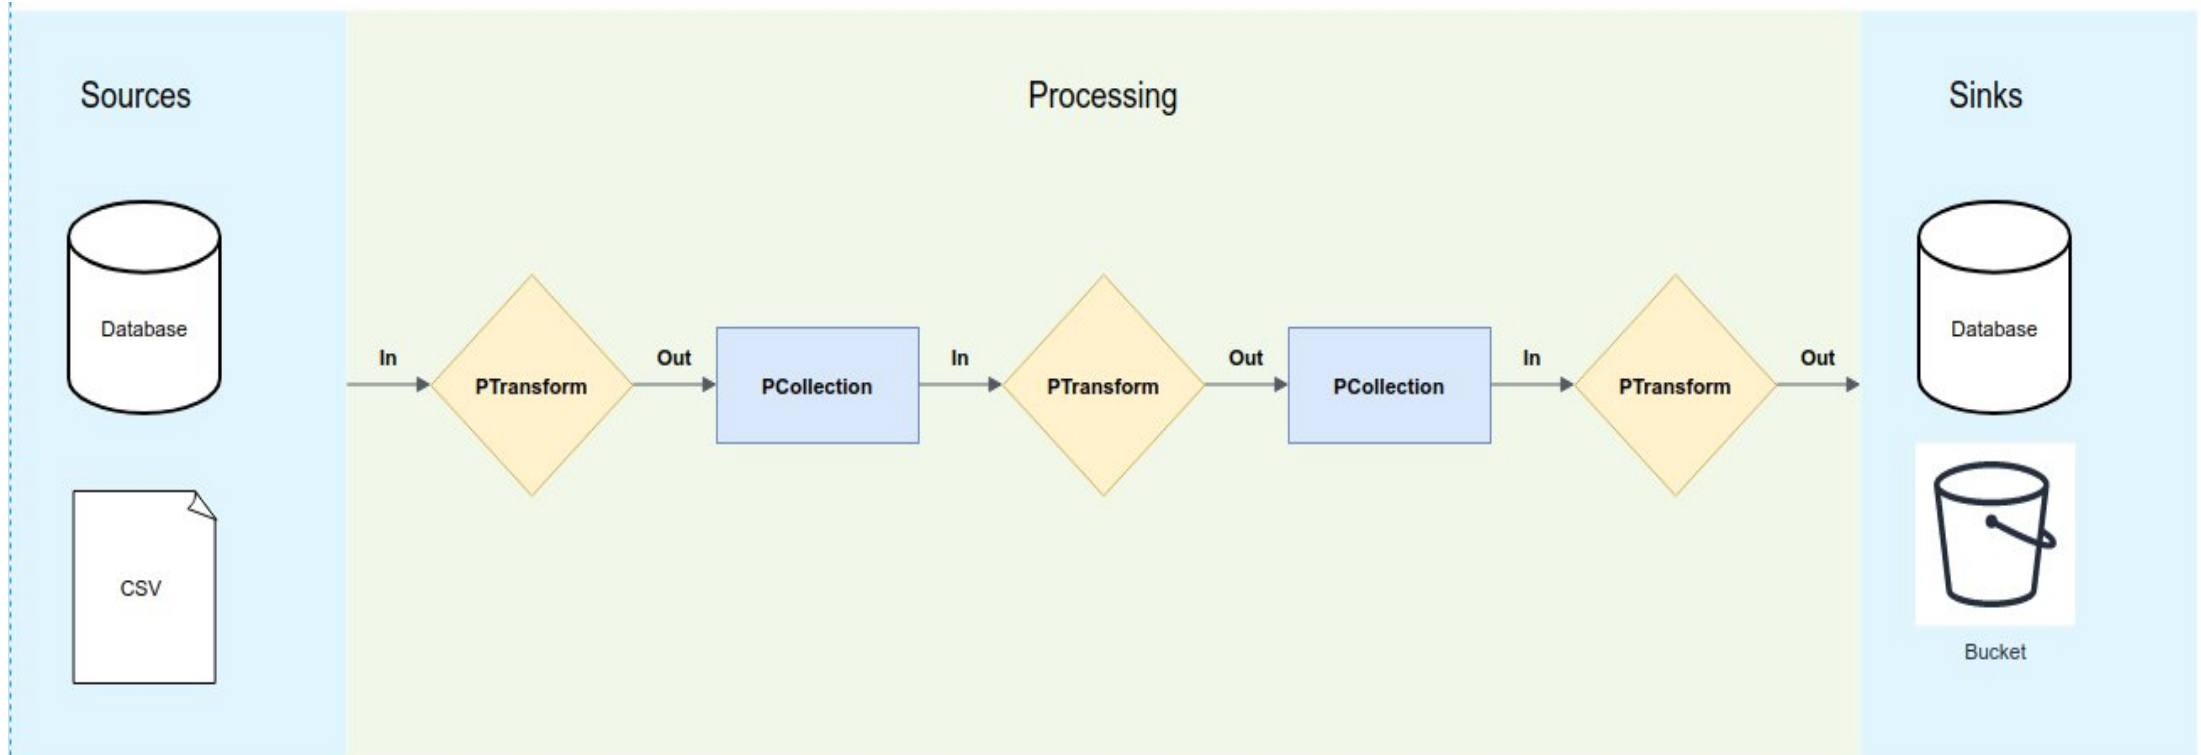
Beam sdk supported languages.



beam

A common ETL flow

Understanding the things mentions before, most of the time our model pipeline looks like.



(Source)

- We need a queue broker like rabbitmq, kafka, AWS SQS or google pub/sub when we are building streaming pipelines.
 - For batch pipelines the source could be a file text, csv or database.

- Streaming common sources



Cloud Pub/Sub




- Batch common sources



(Processing)


- An ETL model we wrote before needs a runner. The supported runners are Google Dataflow, Flink, Spark, samza and more.
- Most known runners supported by Beam



(Sink/Destination)

A common good destination sources are PostgreSQL, S3, cloud storage, MongoDB, Google Bigquery, Amazon RDS.

- Sinks/Destination



Amazon S3



Amazon RDS



PostgreSQL



BigQuery

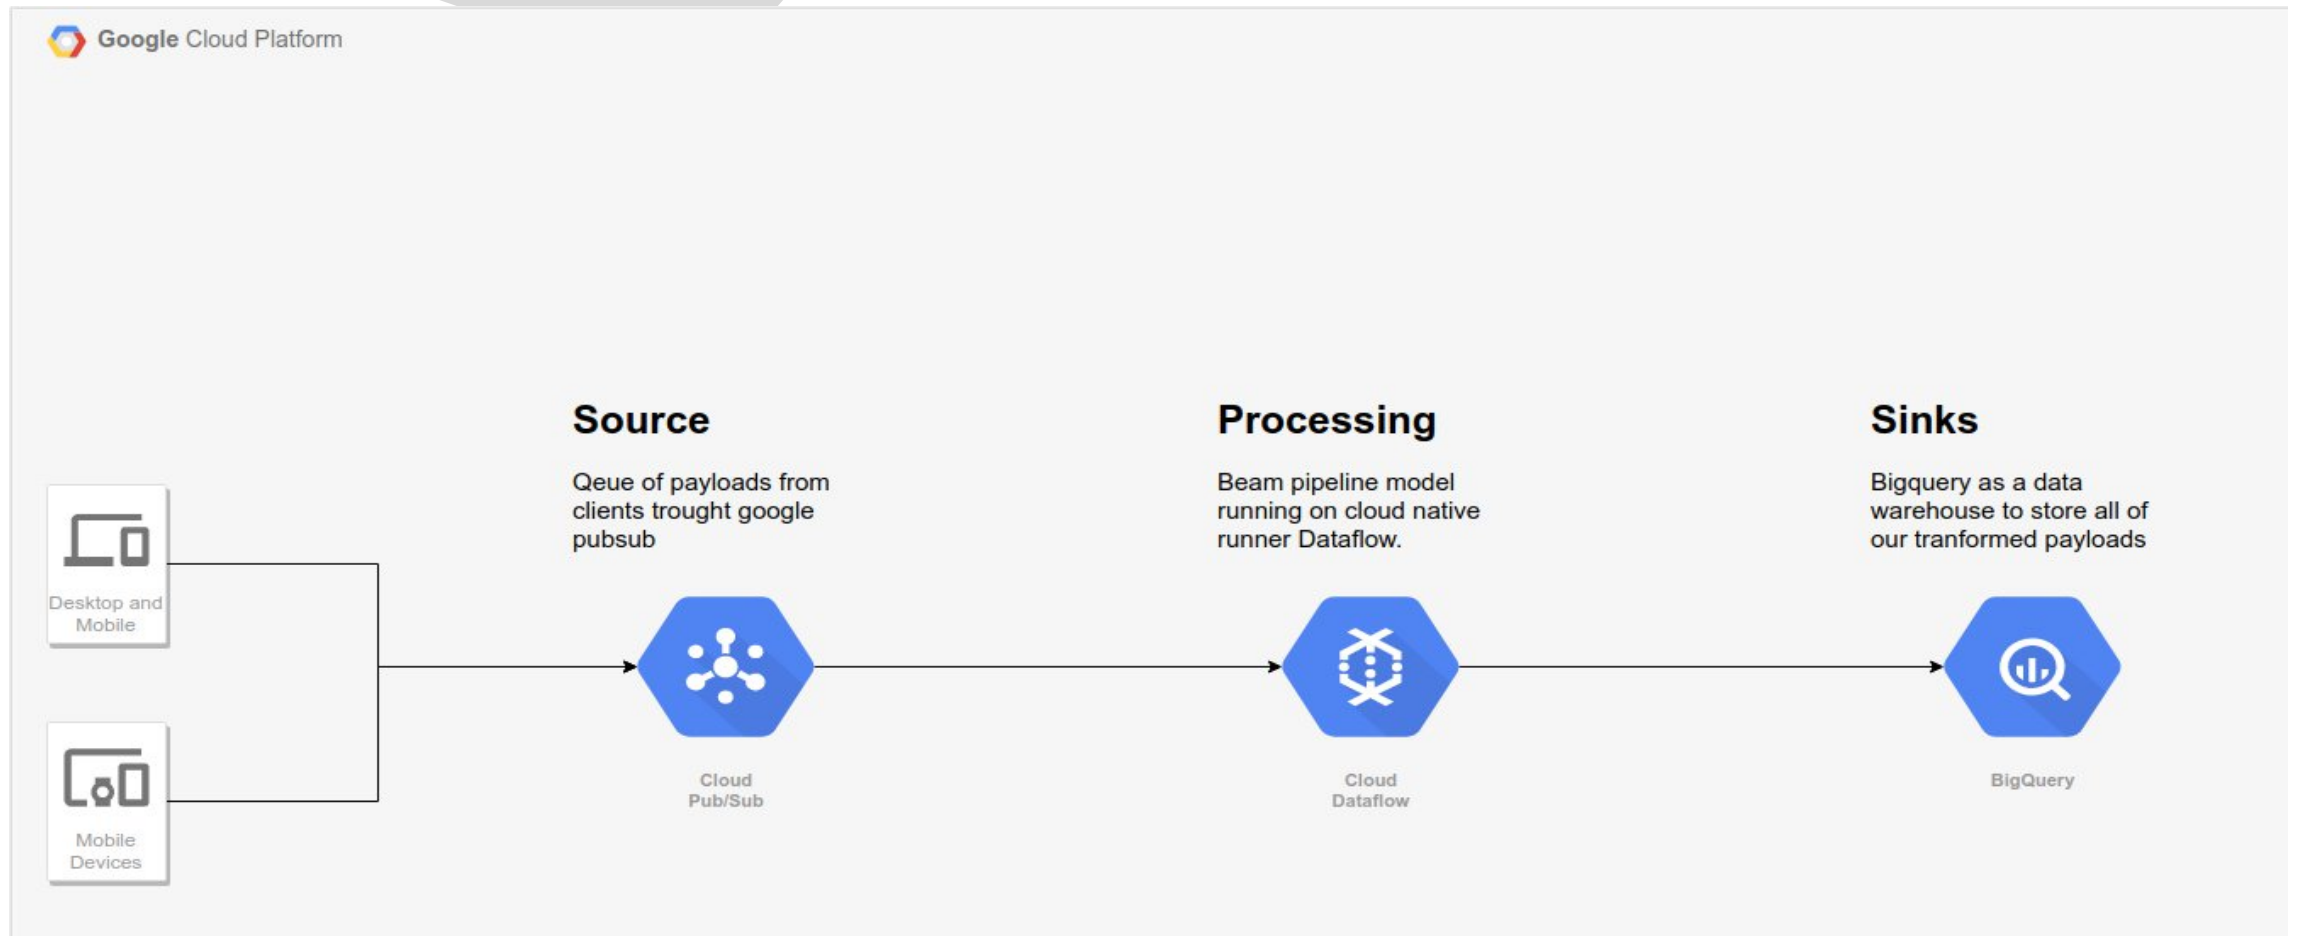


mongoDB®



Architecture for streaming data pipeline.

- Our architecture pipeline looks like



Architecture for streaming data pipeline.

Stages:

- We are going to use google pub/sub because is serverless, auto scalable and not ops config needed (**Source**).
- We are going to use dataflow, because is auto scalable and we don't need a lot of configuration or ops knowledge (**Processing**).
- We are goin to use Google bigquery beacause again here is serverless, auto scalable, without ops config and cheaper only pays for the query executed (**Sink**).

Architecture Performance tips for streaming data pipeline.

- Avoid to make request to another api or external network resource, this can delay or block your pipeline.

```
Events_payload = requests.get('https://api.github.com/events') # tooks 100ms-200ms for each requests
```

```
#If we have 1.000.000 of payloads so 100ms-200ms x 1.000.000, this could be broke the pipeline or have a delay issue
```

Architecture Performance tips for streaming data pipeline.

- Avoid to make validations againsts database or any other resource outside of the logic pipe.

- `user = UserModel.find(id=1234)`
- `If user.active:`
- `#Do something here`
- `Elif user.name == "Joe":`
- `# Another logic here`
-

request a database with queries and validate logic that depends of external resources like this is not recommended for streaming pipelines. The database could have a heavy load and the pipeline can be broke or stop due database response.

Architecture for streaming data pipeline.

- This is our beam processing model for demo



```
b'eyJ1YW11IjogIk1pY2hhZWwgQmVja2VyIiwgImNvbXBhbnkiOiAiSm9obnNvbiBhbmQgU29ucyIsICJtc2ciOiAiU2hvdWxkIGNyaW1lIHByb2R1Y2UgcFpc2Ugc3RhdGVtZW50IG5pY2UgcHJvZHVjdCBhaXIuIiwgInJlbW90ZV9pcCI6ICJyMDEuMjM0LjE0MC4xNTkiLCJlY2VudCI6ICJpcGVyYS84LjUxLihYMTE7IEExpbnV4IHg4Nl82NDsgdGwtUEgpfFByZXN0by8yLjkuMTc2IFZlcnNpb24vMTEuMDA0LCAiZGF0ZSI6ICJyMDIxLTAwLTAwIn0='
```

```
b'{"name": "Michael Becker", "company": "Johnson and Sons", "msg": "Should crime produce raise statement nice product air.", "remote_ip": "201.234.140.159", "user_agent": "Opera/8.51.(X11; Linux x86_64; tl-PH) Presto/2.9.176 Version/11.00", "date": "2021-10-04"}'
```

```
{'name': 'Michael Becker', 'company': 'Johnson and Sons', 'msg': 'Should crime produce raise statement nice product air.', 'remote_ip': '201.234.140.159', 'user_agent': 'Opera/8.51.(X11; Linux x86_64; tl-PH) Presto/2.9.176 Version/11.00', 'date': '2021-10-04', 'ip_info': {'continent': 'South America', 'country': 'Argentina', 'city': 'Buenos Aires'}}
```

```
{'name': 'Michael Becker', 'company': 'Johnson and Sons', 'msg': 'Should crime produce raise statement nice product air.', 'remote_ip': '201.234.140.159', 'user_agent': 'Opera/8.51.(X11; Linux x86_64; tl-PH) Presto/2.9.176 Version/11.00', 'date': '2021-10-04', 'ip_info': {'continent': 'South America', 'country': 'Argentina', 'city': 'Buenos Aires'}, 'device': {'device': 'PC', 'os': 'Linux', 'browser': 'Opera'}}
```

Query complete (0.4 sec elapsed, 0 B processed)

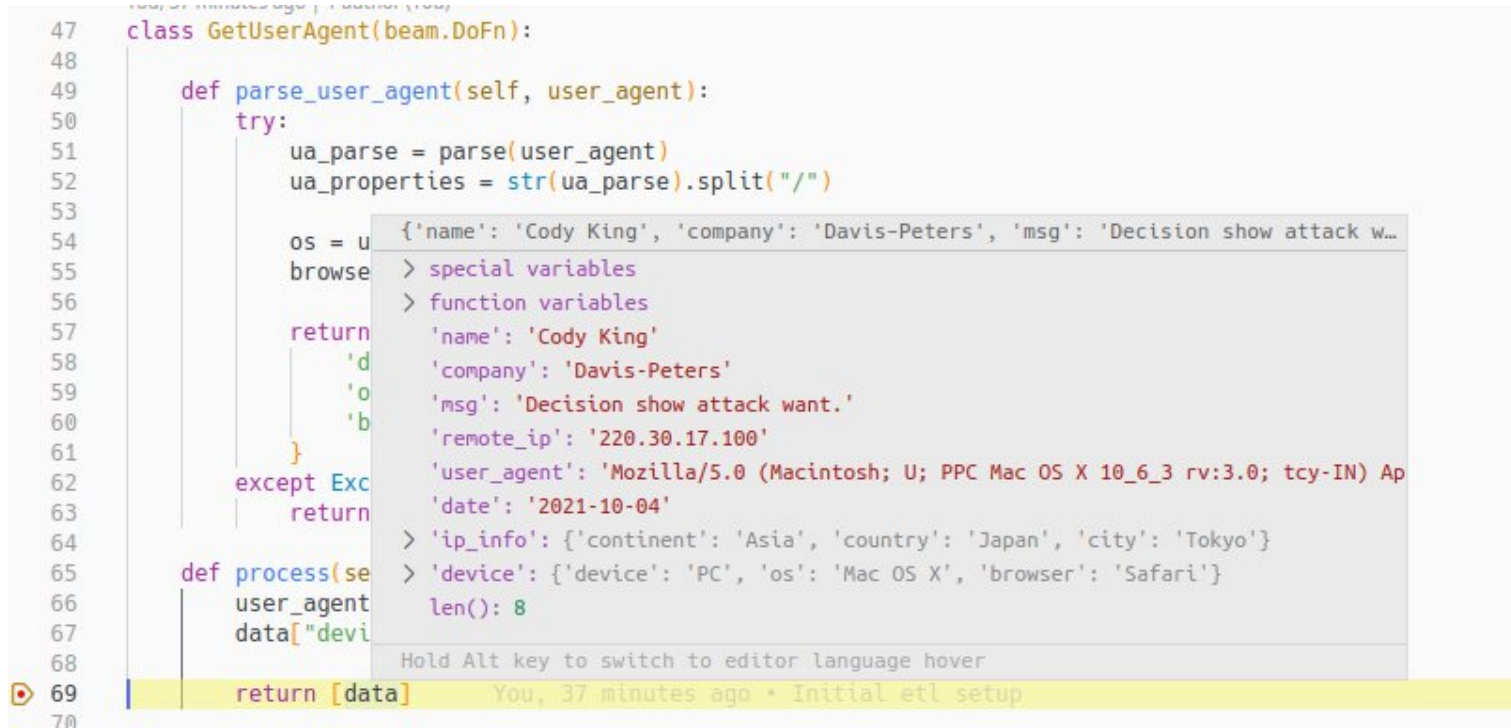
Job Information Results JSON Execution details

Row	name	company	msg	remote_ip	user_agent	date	ip_info.continent	ip_info.country	ip_info.city	device.device	device.os	device.browser
1	Michael Becker	Johnson and Sons	Should crime produce raise statement nice product air.	201.234.140.159	Opera/8.51.(X11; Linux x86_64; tl-PH) Presto/2.9.176 Version/11.00	2021-10-04	South America	Argentina	Buenos Aires	PC	Linux	Opera

Tips and tricks debugging streaming pipeline locally

Apache Beam provide us a DirectRunner which we can use to debug the pipeline in our favorite IDE. (Demo Live)

```
47 class GetUserAgent(beam.DoFn):
48
49     def parse_user_agent(self, user_agent):
50         try:
51             ua_parse = parse(user_agent)
52             ua_properties = str(ua_parse).split("/")
53
54             os = u
55             browse
56
57             return
58             'd
59             'o
60             'b
61         }
62     except Exc
63     return
64
65     def process(se
66     user_agent
67     data["devi
68
69     return [data]
```



Hold Alt key to switch to editor language hover

You, 37 minutes ago • Initial etl setup


Deploying the data pipeline in cloud runner (Demo Live).

Topic details

Topic name `projects/demos-304101/topics/demo-beam`

Publish message request count

Requests/sec



The screenshot shows the BigQuery interface. At the top, there are tabs for RUN, SAVE, SCHEDULE, and MORE. Below the tabs is a query editor with the following SQL query: `SELECT * FROM `demos-304101.demobeam.users_data` WHERE date = "2021-10-04" LIMIT 1000`. Below the query editor, there is a section for Query results. It shows the processing location as US. There are links for SAVE RESULTS and EXPLORE DATA. Below these links, it says "Query complete (0.4 sec elapsed, 0 B processed)". There are tabs for Job information, Results, JSON, and Execution details. The Results tab is selected, showing a table with the following data:

Row	name	company	msg	remote_ip	user_agent	date	ip_info.continent	ip_info.country	ip_info.city	device.device	device.os	device.browser
1	Michael Becker	Johnson and Sons	Should crime produce raise statement nice product air.	201.234.140.159	Opera/8.51.(X11; Linux x86_64; 6-PM) Presto/2.9.176 Version/11.00	2021-10-04	South America	Argentina	Buenos Aires	PC	Linux	Opera